

BIJAY THAKUR

[East Elmhurst, New York](#)

[+1 \(330\) 883-0416](#)

[starbijay143@gmail.com](#)

[LinkedIn](#)

[GitHub](#)

[Portfolio](#)

SUMMARY

AI Engineer specializing in production RAG and multi-agent orchestration. Shipped a live financial QA platform with Supabase auth, Stripe-gated quotas, and Google Cloud Run end-to-end across OpenAI, Gemini, LangGraph, LlamaIndex, ChromaDB, and FastAPI.

EDUCATION

Queens College, The City University of New York

Expected: Dec 2026

*Bachelor of Science in Computer Science, **GPA: 3.64***

PROJECTS

SEC Insight AI — Production Financial QA System | Live Demo

April 2026 - June 2026

- Engineered a citation-grounded RAG pipeline over Apple's 2025 annual report with OpenAI embeddings, ChromaDB, BM25 hybrid retrieval, RRF, and cross-encoder reranking - achieving 0.9901 RAGAS faithfulness and 0.8032 MRR on a 50-question gold eval set.
- Deployed full-stack to production: Dockerized FastAPI backend on Google Cloud Run, Next.js frontend on Vercel, Supabase auth with per-user quota enforcement, tiered rate limiting (30–300 req/hr), answerability gate, and server-side cost guards - with 9-file test suite and CI/CD via GitHub Actions.

LaunchLab AI — Agentic Multi-LLM Product Studio | GitHub

April 2026 - May 2026

- Architected a multi-LLM agentic system using LangGraph, Flask, Gemini, and OpenAI to orchestrate 7 specialized agents - producing market research, product blueprints, Claude Code build prompts, campaign copy, visual prompts, and evaluation reports from a single product brief.
- Implemented manual v1 and LangGraph v2 orchestration modes with identical response shapes, provider fallback logic (live → partial → mock), and structured error handling - enabling direct latency and reliability benchmarking between orchestration approaches.

Code Council — Multi-Agent GenAI Code Analysis | GitHub

Jan 2026 - May 2026

- Built a 6-agent code analysis system (Orchestrator, Explainer, Bug Hunter, Complexity Analyst, Debate rounds, Supervisor) using Express and React/TypeScript, with Explainer/BugHunter/Complexity executing in parallel.
- Implemented Zod schema validation on all LLM outputs, trace ID logging, fallback-safe responses, and a context-aware chat assistant with full analysis context - preventing hallucinated or malformed responses from breaking the UI.

MarianMT Domain Adaptation — Neural Machine Translation Fine-Tuning

May 2026

- Fine-tuned Helsinki-NLP/opus-mt-en-fr using Hugging Face Seq2SeqTrainer, PyTorch, and mixed-precision CUDA on Google Colab across 1K, 10K, and 189K KDE4 training examples - improving English-to-French BLEU from 46.07 to 58.37 (+26.7%) and validating consistent scaling improvement across all dataset sizes.

EXPERIENCE

AI Agent Builder Intern, NeuralSeek

April 2026 - May 2026

Remote

- Engineered multi-step AI agent workflows on the mAIstro no-code platform, iterating on prompt design, reasoning chains, and validation logic to improve reliability of production automation pipelines.
- Systematically stress-tested agent outputs across edge cases, identified failure modes, and produced structured failure documentation that directly informed pipeline refinements across multiple workflows.

Academic Tutor, LaGuardia Community College

Feb 2024 - Jun 2024

Long Island City, NY

- Tutored students 1-on-1 in C++, debugging, and algorithmic problem-solving; developed structured practice modules to accelerate mastery of core CS fundamentals.

TECHNICAL SKILLS

- Languages:** Python, TypeScript, JavaScript, Java, C++, SQL
- AI & ML:** RAG Pipelines, Hybrid Retrieval (BM25 + RRF), Cross-Encoder Reranking, Multi-Agent Orchestration, LLM Fine-Tuning, LangGraph, LlamaIndex, Hugging Face Transformers, PyTorch
- Backend/DevOps:** FastAPI, Flask, Express, REST APIs, Docker, Google Cloud Run, GitHub Actions